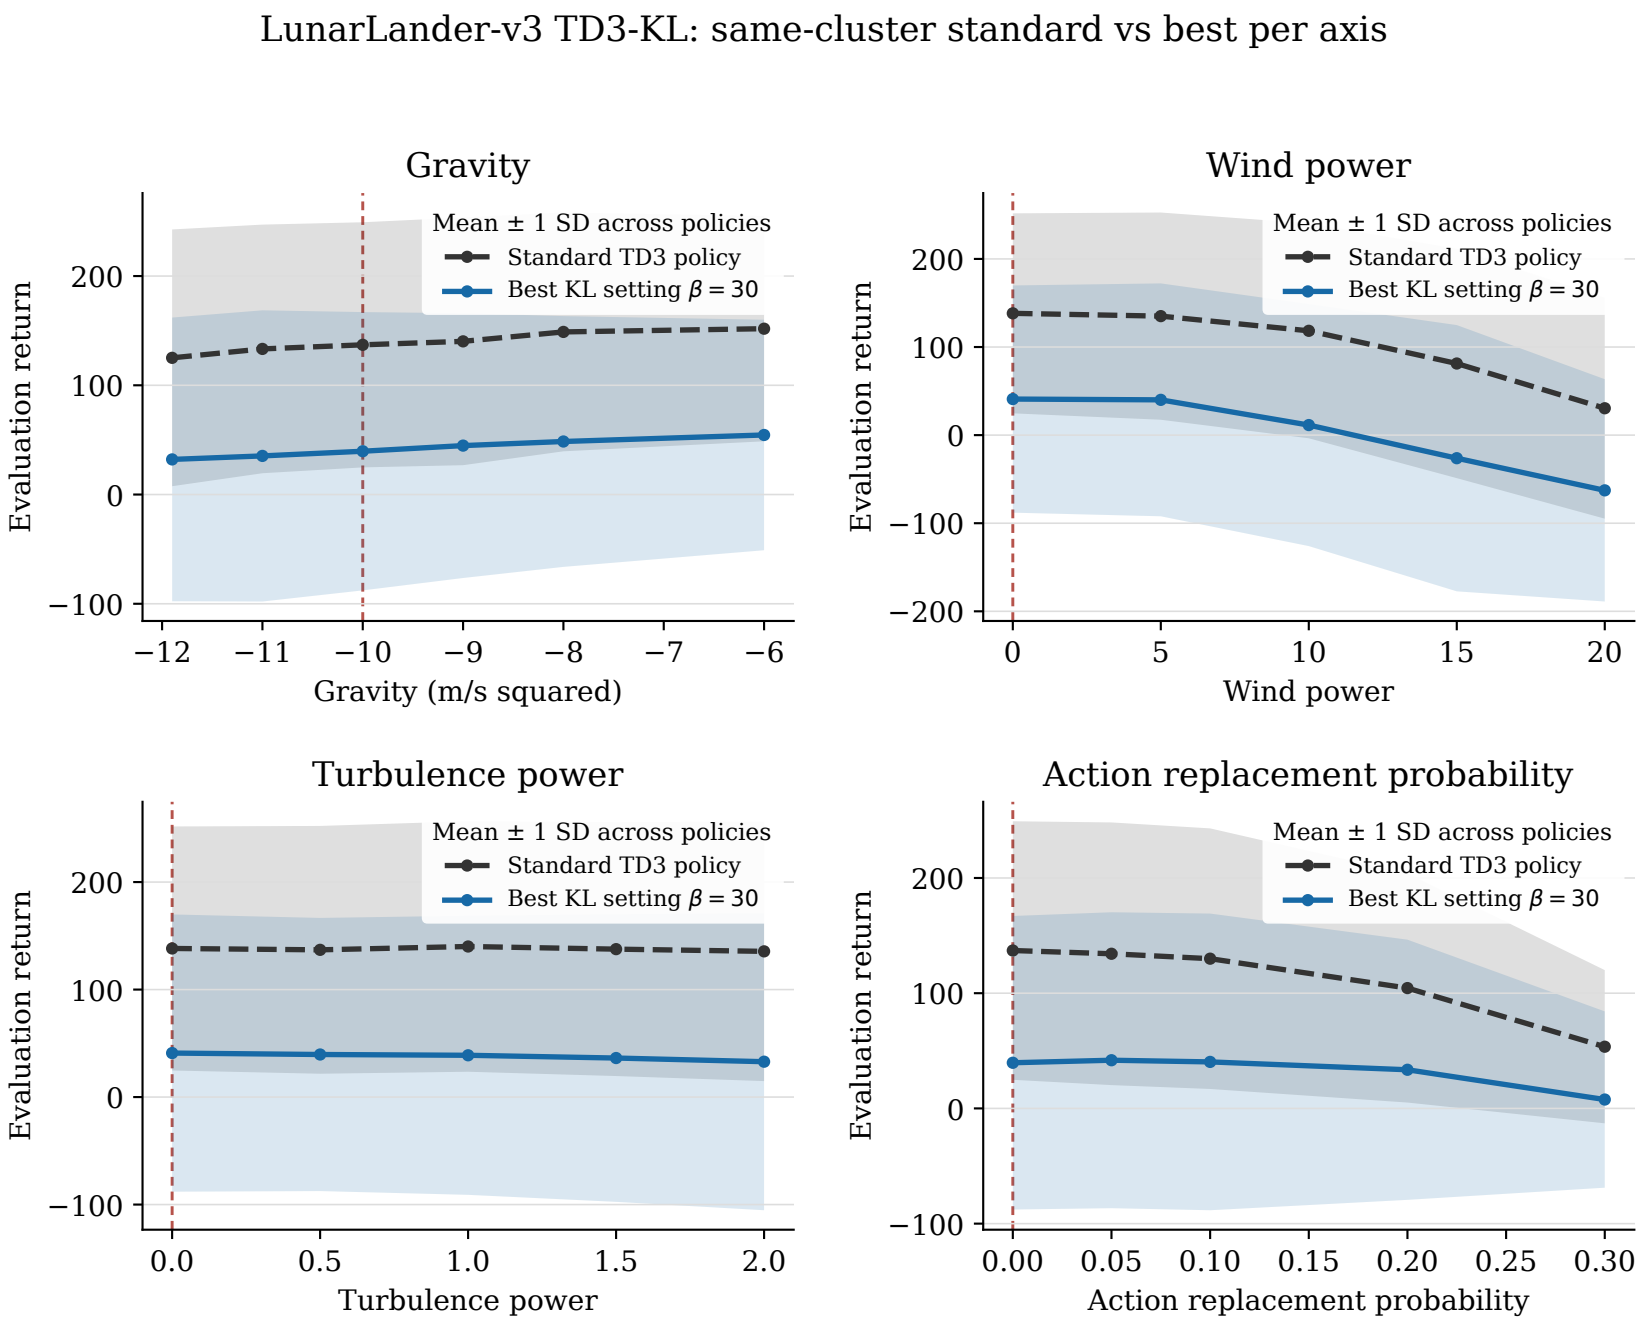
LunarLander-v3 TD3-KL: same-cluster standard vs best per axis
Gravity
Wind power
Turbulence power
Action replacement probability
Mean ± 1 SD across policies
Standard TD3 policy
Best KL setting β = 30
Evaluation return
Gravity (m/s squared)
Wind power
Turbulence power
Action replacement probability
30 training seeds | Mean +/- sample SD across policy means | Ex-post per-axis selection