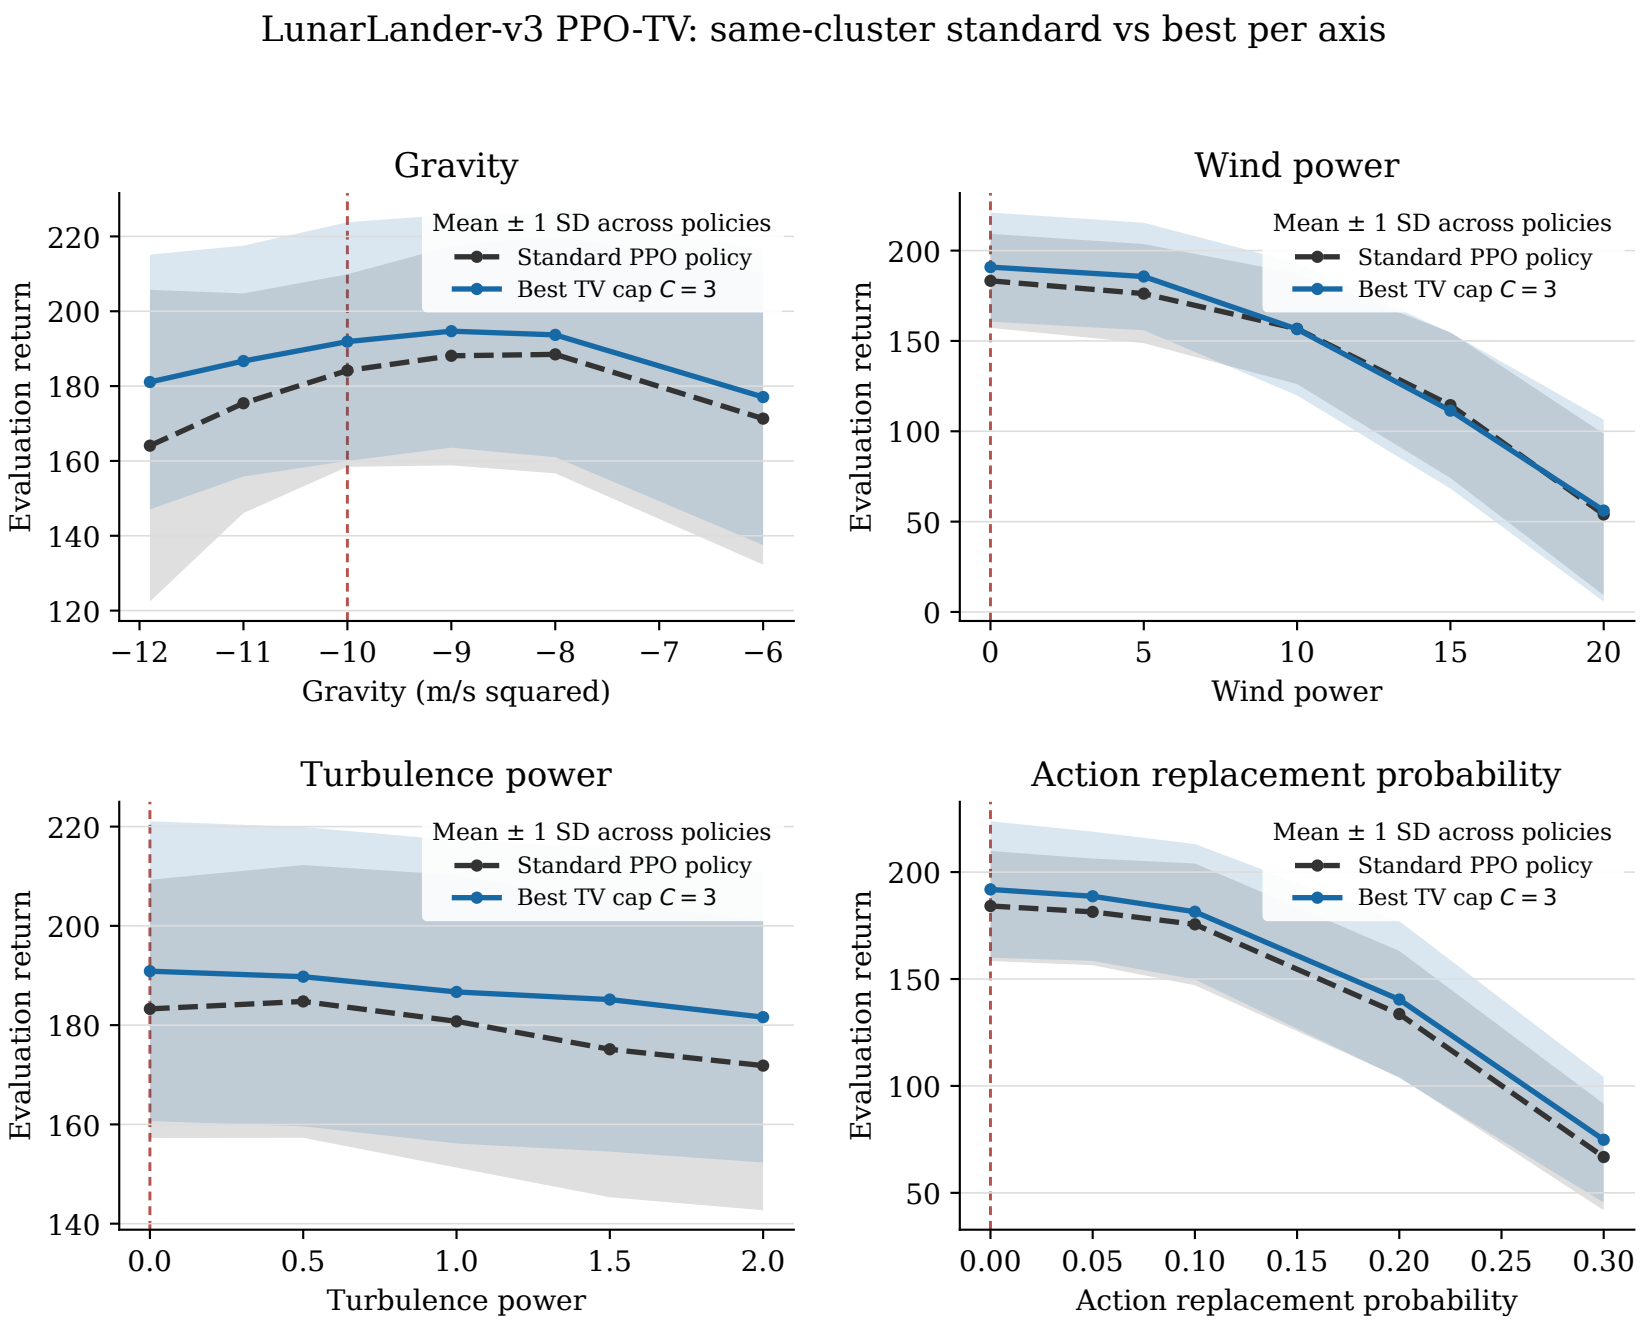
LunarLander-v3 PPO-TV: same-cluster standard vs best per axis
Gravity
Wind power
Turbulence power
Action replacement probability
Mean ± 1 SD across policies
Standard PPO policy
Best TV cap $C = 3$
Evaluation return
Gravity (m/s squared)
Wind power
Turbulence power
Action replacement probability
30 training seeds | Mean +/- sample SD across policy means | Ex-post per-axis selection