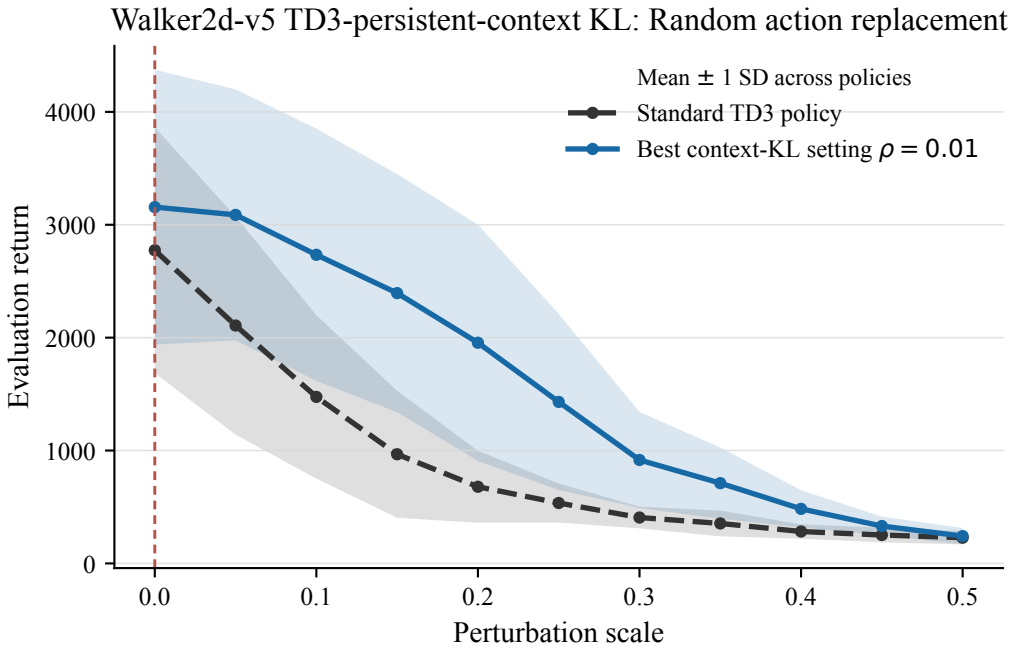
Walker2d-v5 TD3-persistent-context KL: Random action replacement
Mean ± 1 SD across policies
Standard TD3 policy
Best context-KL setting ρ = 0.01
Evaluation return
4000
3000
2000
1000
0
0.0
0.1
0.2
0.3
0.4
0.5
Perturbation scale