Suite-level effects: a reliability gain can coexist with a return loss

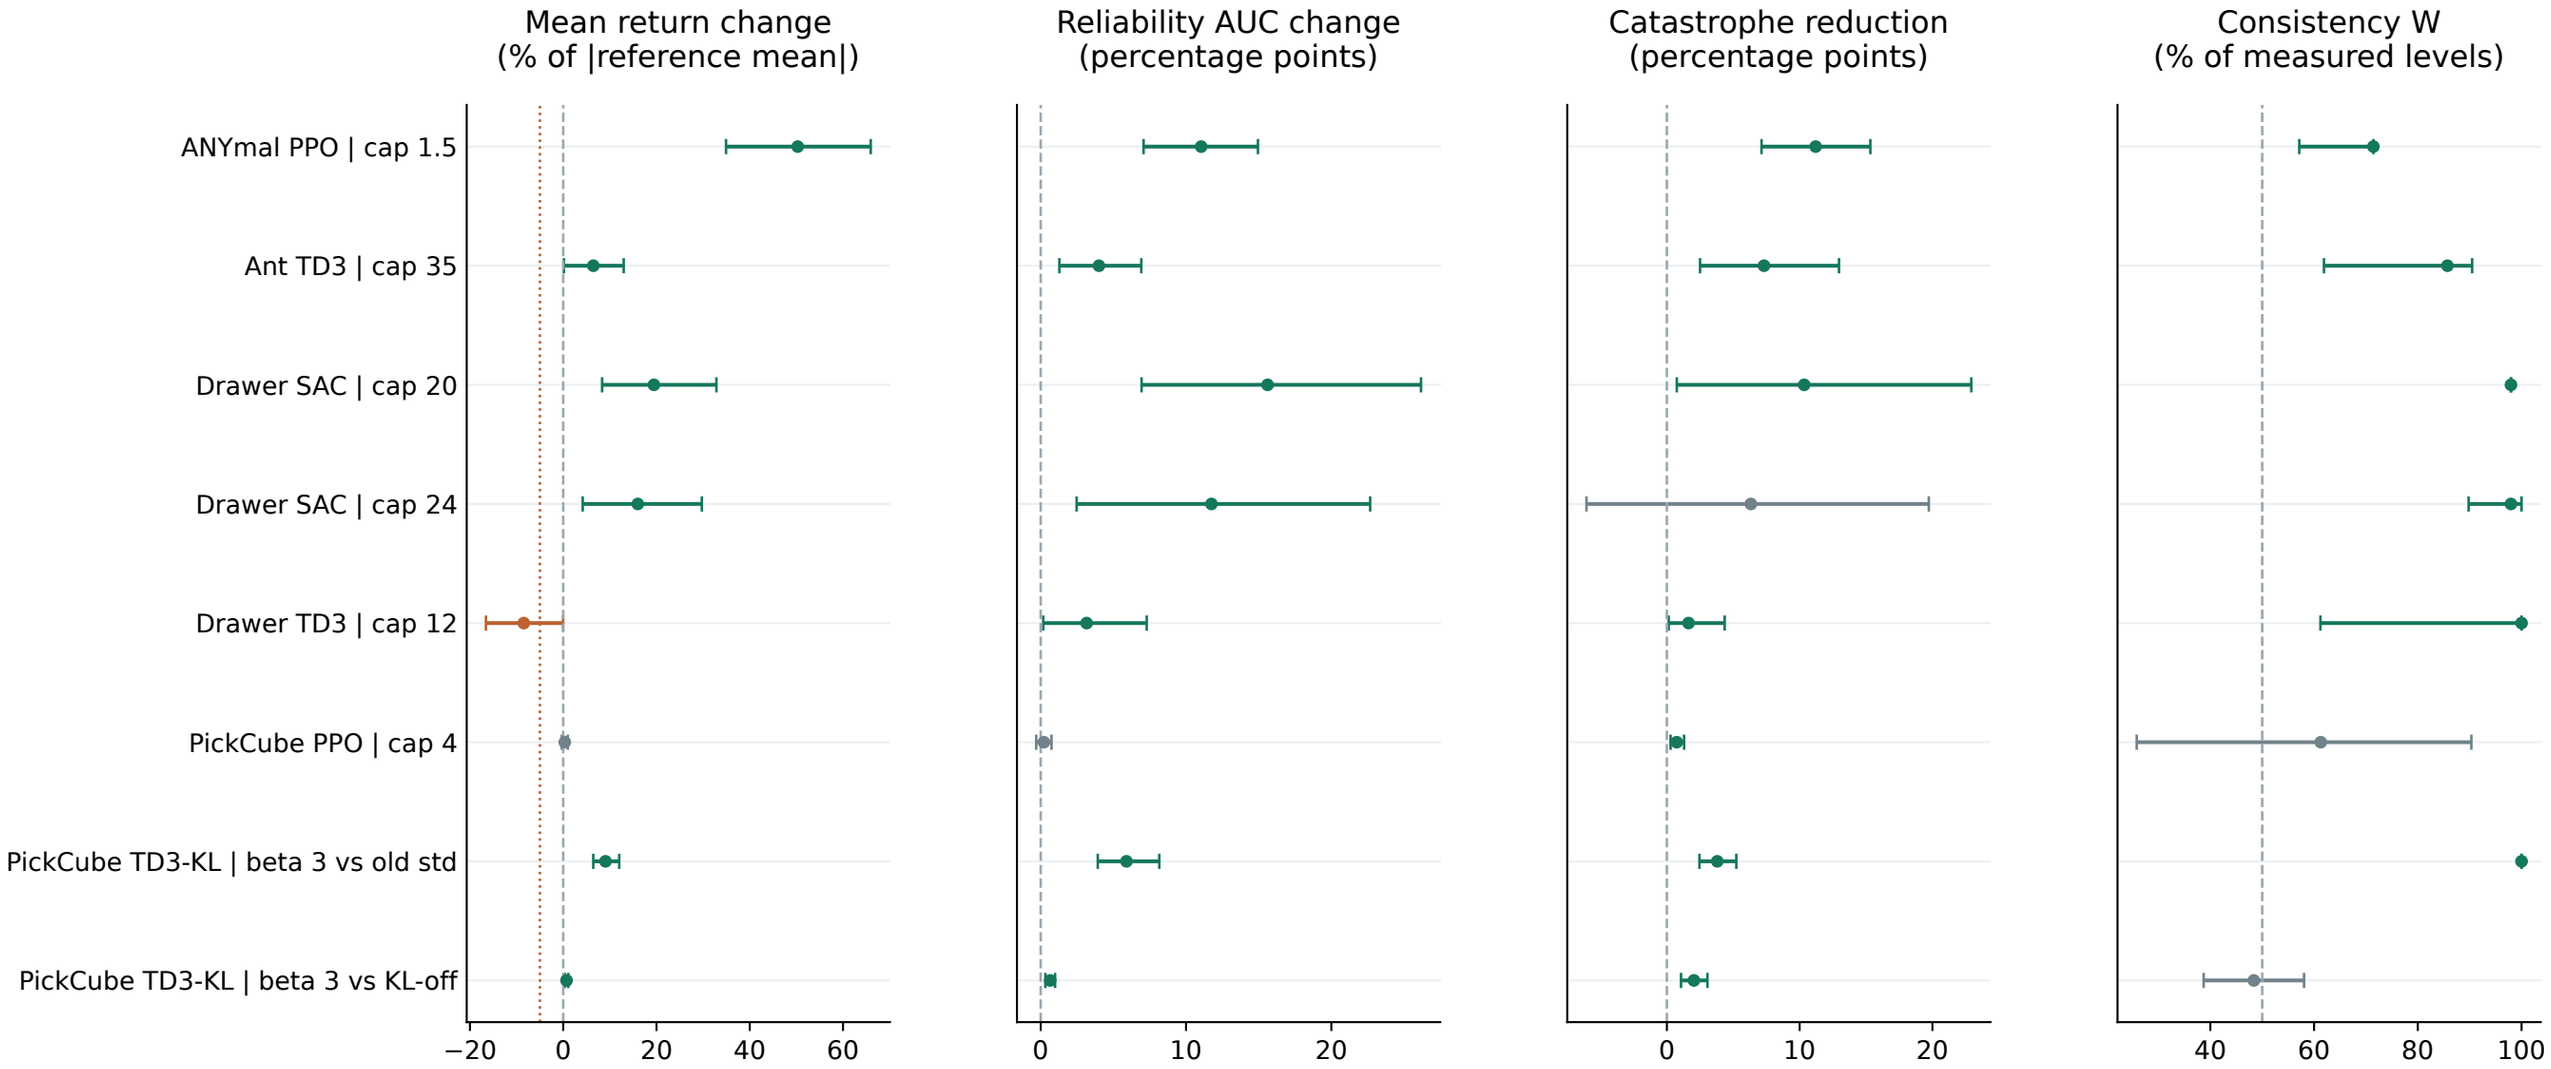

30 policies per method; 10,000 whole-policy bootstrap draws. Design-appropriate pairing, cross-batch old-standard comparisons independent.
Dotted return line: -5% noninferiority margin. Mean changes are normalized for display only. Pointwise intervals; no multiplicity correction.